



# Traffic filtering at scale on Linux

François Serman <[francois.serman@corp.ovh.com](mailto:francois.serman@corp.ovh.com)>

Pass The Salt 2018

- Introduction
- (past) BPF
- (present) eBPF
- Let's play with BPF!
- Performance analysis
- Summary and conclusion

# Introduction

# whoami

```
fserman@ovh $ groups
```

```
dev vac
```

```
fserman@ovh $ uptime | awk '{ print $2, $3, $4 }'
```

```
up 435 days,
```


```
fser@home $ groups
```

```
clx, lautre.net, hexpresso
```

# Back to the presentations


- ▶ Traffic filtering:
  - ▶ Obviously: classify packets we want to keep, drop the rest ;
  - ▶ Achieved using (e)BPF.
- ▶ at scale:
  - ▶ Tenth of gigabits per seconds ;
  - ▶ Millions of packets per seconds ;
  - ▶ We'll see how to generate such traffic ;
  - ▶ but also how to mitigate it (XDP).
- ▶ on Linux:
  - ▶ Using recent ( $> 4.8$ ) kernel facilities.

# Networking 101



# Breadcrumb

Top amplification attack on Memcached (UDP 11211) : 1.3Tbps.  
(For the record: MIRAI was 1Tbps)



The amplification attack aiming Memcached in march 2018.

(past) BPF



## 199[23] : Steven McCanne &amp; Van Jacobson at Berkeley

## The BSD Packet Filter: A New Architecture for User-level Packet Capture

*Steven McCanne & Van Jacobson – Lawrence Berkeley Laboratory<sup>1</sup>*

### ABSTRACT

Many versions of Unix provide facilities for user-level packet capture, making possible the use of general purpose workstations for network monitoring. Because network monitors run as user-level processes, packets must be copied across the kernel/user-space protection boundary. This copying can be minimized by deploying a kernel agent called a *packet filter*, which discards unwanted packets as early as possible. The original Unix packet filter was designed around a stack-based filter evaluator that performs sub-optimally on current RISC CPUs. The BSD Packet Filter (BPF) uses a new, register-based filter evaluator that is up to 20 times faster than the original design. BPF also uses a straightforward buffering strategy that makes its overall performance up to 100 times faster than Sun's NIT running on the same hardware.

Provide a way to filter packets and avoid useless packets copies (kernel to user).

# Main concepts

- ▶ [Efficient] Kernel architecture for packet capture;
  - ▶ Discard unwanted packets as early as possible;
  - ▶ Packet data references should be minimised;
  - ▶ Decoding an instruction ~ single C switch statement;
  - ▶ Abstract machine registers should reside in physical one;
- ▶ Protocol independent: no modification to the kernel to support a new protocol;
- ▶ General: instruction set should be rich enough to handle unforeseen uses;

# BPF is a virtual machine

What is a virtual machine?

- ▶ Abstract computing machine;
- ▶ Has its own instruction-set, registers, memory representation;
- ▶ Cannot run directly on actual hardware:
- ▶ Hence need a VM loader and interpreter or compiler.

# The BPF virtual machine

All values are 32 bits (instructions / data)

Fixed-length instructions:

- ▶ **Load** data to registers;
- ▶ **Store** data to memory;
- ▶ **ALU instructions** arithmetic or logic operations;
- ▶ **Branch instructions** alter the control-flow based on a test;
- ▶ **Return instructions** terminate the filter;
- ▶ **(Misc operations)**

# Usage

Most famous use case:

- ▶ **tcpdump** (via **libpcap**).
- ▶ **cls\_bpf** (TC classifier for shaping)
- ▶ **xt\_bpf** (iptables module).

Please tcpdump, show us all **UDP** packets towards **memcached**.

```
# tcpdump -p -d 'ip and udp and dst port 11211'
```

Notice the difference with/without *«ip and»*

# Under the hood

```
# tcpdump -p -d 'ip and udp and dst port 11211'
(000) ldh      [12]
(001) jeq      #0x800          jt 2      jf 10
(002) ldb      [23]
(003) jeq      #0x11          jt 4      jf 10
(004) ldh      [20]
(005) jset     #0x1fff          jt 10     jf 6
(006) ldx      4*([14]&0xf)
(007) ldh      [x + 16]
(008) jeq      #0x2bcb          jt 9      jf 10
(009) ret      #262144
(010) ret      #0
```

## Decrypting the output

- ▶ (000) ldh [12]  
Load half-word from packet at offset 12 (EtherType)

## Decrypting the output

- ▶ (000) `ldh [12]`  
Load half-word from packet at offset 12 (EtherType)
- ▶ (001) `jeq #0x800 jt 2 jf 10`  
If equals 0x800 (EtherType IPv4). If true, go to 2, else to 10.



## Decrypting the output

- ▶ (000) `ldh [12]`  
Load half-word from packet at offset 12 (EtherType)
- ▶ (001) `jeq #0x800 jt 2 jf 10`  
If equals 0x800 (EtherType IPv4). If true, go to 2, else to 10.
- ▶ (002) `ldb [23]`  
Load double-word at offset 23 (Protocol field in IPv4 header)

# Decrypting the output

- ▶ (000) `ldh [12]`  
Load half-word from packet at offset 12 (EtherType)
- ▶ (001) `jeq #0x800 jt 2 jf 10`  
If equals 0x800 (EtherType IPv4). If true, go to 2, else to 10.
- ▶ (002) `ldb [23]`  
Load double-word at offset 23 (Protocol field in IPv4 header)
- ▶ (003) `jeq #0x11 jt 4 jf 10`  
If proto is UDP, continue to 4, else go to 10

# Decrypting the output


- ▶ (000) ldh [12]  
Load half-word from packet at offset 12 (EtherType)
- ▶ (001) jeq #0x800 jt 2 jf 10  
If equals 0x800 (EtherType IPv4). If true, go to 2, else to 10.
- ▶ (002) ldb [23]  
Load double-word at offset 23 (Protocol field in IPv4 header)
- ▶ (003) jeq #0x11 jt 4 jf 10  
If proto is UDP, continue to 4, else go to 10
- ▶ (007) ldh [x + 16]  
Load UDP Dest port

# Decrypting the output

- ▶ (000) ldh [12]  
Load half-word from packet at offset 12 (EtherType)
- ▶ (001) jeq #0x800 jt 2 jf 10  
If equals 0x800 (EtherType IPv4). If true, go to 2, else to 10.
- ▶ (002) ldb [23]  
Load double-word at offset 23 (Protocol field in IPv4 header)
- ▶ (003) jeq #0x11 jt 4 jf 10  
If proto is UDP, continue to 4, else go to 10
- ▶ (007) ldh [x + 16]  
Load UDP Dest port
- ▶ (008) jeq #0x2bcb jt 9 jf 10  
If dest port == 11211 (0x2bcb), go to 9, else go to 10

# Visualization

```
tcpdump -p -d 'ip and udp and dst port 11211'
```



(present) eBPF

# Improvements (~ 2013)

From Documentation/networking/filter.txt:

- ▶ Registers:
  - ▶ Increase number of registers from 2 to 10;
  - ▶ 64 bits formats;
  - ▶ ABI mapped on the underlying architecture;
- ▶ Operations in 64 bits;
- ▶ Conditionnal jt/jf replaced with jt/fall-through;
- ▶ BPF calls;
- ▶ Maps

# eBPF today

- ▶ the old BPF is referred to as classic BPF (cBPF);
- ▶ eBPF is the new BPF!
- ▶ No longer limited to packet filtering:
  - ▶ tracing (kprobes);
  - ▶ security (seccomp);
  - ▶ ...



# eBPF today

- ▶ BPF is very suitable for *JIT* (Just In Time compilation):
  - ▶ Virtual registers already map the physicals one;
  - ▶ Only have to issue the proper instruction;
  - ▶ Available for x86\_64, arm64, ppc64, s390x, mips64, sparc64 and arm;
  - ▶ 1 C switch statement became 1 instruction.
- ▶ BPF bytecode is **verified** before loading in the kernel.
- ▶ Hardened JIT available.

```
# echo 1 > /proc/sys/net/core/bpf_jit_enable
```

# eBPF verifier

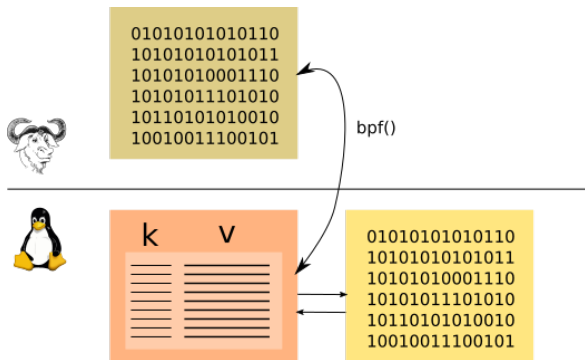
Provides a verdict whether the bytecode is safe to run:

- ▶ a BPF program must **always** terminate:
  - ▶ size-bounded (max 4096 instr);
  - ▶ Loop detections (CFG validation);
- ▶ a BPF program must be safe:
  - ▶ detecting out of range jumps
  - ▶ detecting out of bounds r/w
  - ▶ context-aware: verifying helper function call's arguments
  - ▶ ...

Refere to *kernel/bpf/verifier.c*.

# eBPF Maps (1/3)

Generic storage facility for sharing data between kernel and userspace.



Interact via *bpf()* syscall (lookup/update/delete).  
Helpers available on *tools/lib/bpf/bpf.h*.

## eBPF Maps (2/3)

Defined by:

- ▶ types (as of 4.18 19 types):
  - ▶ **Arrays** *BPF\_MAP\_TYPE\_ARRAY* (+ PERCPU);
  - ▶ **Hashes** *BPF\_MAP\_TYPE\_HASH* (+PERCPU);
  - ▶ **LRU** *BPF\_MAP\_TYPE\_LRU\_HASH* (+PERCPU);
  - ▶ **LPM** *BPF\_MAP\_TYPE\_LPM\_TRIE*;
- ▶ max number of elements
- ▶ key size in bytes
- ▶ value size in bytes

Let's play with BPF!

# In kernel tools

Have a look on *[samples/bpf](#)*:

- ▶ `bpf_asm` a minimal cBPF assembler;
- ▶ `bpf_dbg` a small debugger for cBPF programs;
- ▶ `bpftool` a generic tool to interact with eBPF programs:
  - ▶ show dump load pin programs
  - ▶ show create pin update delete maps
  - ▶ ...

# BPF Compiler Collection (BCC)

Quoting their README:

- ▶ “Toolkit for creating efficient kernel tracking and manipulation programs [...]”
- ▶ “it makes use of extended BPF”.

For us:

- ▶ Provides a way to load BPF code (not only for networking)
- ▶ Collection of BPF programs (traces, perf. . . )
- ▶ Python API

# Demo time

## Collect statistics on running memcached.

- ▶ One party generates memcached requests (randomly);
- ▶ The other party has two parts:
  - ▶ kernel part: parses the protocol, extracts the request's keyword, and updates counters;
  - ▶ userspace part: periodically displays the counters.

## Memcached commands:

add append cas decr delete flush\_all get gets incr prepend replace stats

```
$ wc -l *  
    30 flood.py  
   188 xdp_memcached.c  
   144 xdp_memcached.py
```



# Performance analysis

# Some numbers

- ▶ Achieving high bandwidth is “easy”
- ▶ Handling lots of packets is harder:
  - ▶ For 64bytes pkts ( $\sim 80$  on the wire)
    - ▶ 10Gbps : 14.8Mpps
    - ▶ 25Gbps : 37.0Mpps
    - ▶ 50Gbps : 74.0Mpps
    - ▶ 100Gbps: 148.0Mpps
  - ▶ For 1500 bytes pkts:
    - ▶ 10Gbps : 820Kpps
    - ▶ 25Gbps :  $\sim 2$ Mpps
    - ▶ 50Gbps :  $\sim 4.1$ Mpps
    - ▶ 100Gbps:  $\sim 8.2$ Mpps

# Experimental setup

- ▶ Two servers : one sender and one receiver
  - ▶ 2 \* Intel(R) Xeon(R) Gold 6134 CPU @ 3.20GHz (8c/16t)
  - ▶ 12 \* 8Gb (= 96Gb) DDR4
  - ▶ Mellanox MT27700 (50Gbps ConnectX-4)
  - ▶ Linux v4.15
- ▶ back to back (no switch was harmed for this presentation)

## Objectives:

- ▶ Increase PPS
- ▶ Reduce packet loss

# Produce modern graphs

Install the following packages:

- ▶ InfluxDB
- ▶ Telegraf
- ▶ Grafana

Import dashboard **928**.

Done.

# State of the art Yolo devops

```
# wget https://dl.influxdata.com/influxdb/releases/ \
influxdb_1.1.1_amd64.deb
# wget https://dl.influxdata.com/telegraf/releases/ \
telegraf_1.1.2_amd64.deb
# wget https://s3-us-west-2.amazonaws.com/ \
grafana-releases/release/grafana_5.1.4_amd64.deb

# dpkg -i *.deb

# sed -i 's/^# \([\\[inputs\\.net\\]\\)/\\1/' \
/etc/telegraf/telegraf.conf

# systemctl start {influxdb,telegraf,grafana-server}.service
```

# Generating traffic

We'll cover several methods to generate traffic. You'll have to guess the rate (in pps) for each:

- ▶ `while true; do nc ... ; done`
- ▶ `python flood.py`
- ▶ `scapy`
- ▶ `tcpreplay`
- ▶ C threaded program
- ▶ kernel's `pktgen`
- ▶ DPDK's `pktgen`

## netcat (code)

```
while true ; do
  ( echo 'Hello, world!' |
    nc -w 1 -u 10.0.1.2 $((RANDOM %65534)) & )
done
```

## netcat (outcome)





# python (code)

```
import socket

UDP_IP, UDP_PORT = "10.0.1.2", 5005
MESSAGE = "Hello, World!"

if len(sys.argv) == 2:
    UDP_PORT = int(sys.argv[1])

sock = socket.socket(socket.AF_INET, # Internet
                     socket.SOCK_DGRAM) # UDP

while True:
    sock.sendto(MESSAGE, (UDP_IP, UDP_PORT))
```

## python (outcome)



## python (multiple processes)

```
for i in {4000..4032} ; do  
    ( python flood.py ${i} & )  
done
```

## python multiple processes (outcome)



## scapy (code)

```
send(IP(dst="10.0.1.2")/UDP(dport=123), loop=100000)
```

## scapy (outcome)



## tcpreplay (code)

```
>>> wrpcap("/tmp/batch.pcap",  
          Ether(dst="7c:fe:90:57:ab:c8")  
          / IP(src="10.0.1.1",dst="10.0.1.2")  
          / UDP(dport=123) * 1000)  
# tcpreplay -i enp134s0f0 --loop 5000000 -tK /tmp/batch.pcap
```

Where *-t* stands for “topspeed” and *k* ...

## tcpreplay (outcome)





## C threaded program (code)

- ▶ <https://github.com/vbooter/DDoS-Scripts/blob/master/UDP.c>
- ▶ (minor modification)

```
# ./UDP 10.0.1.2 4242 0 64 32
```

- ▶ 0 is the throttle
- ▶ 64 the packet size
- ▶ 32 the number of threads

# C threaded program (outcome)



## kernel's pktgen (config)


```
# cd ~/linux/sample/pktgen
# export PGDEV=/proc/net/pktgen/enp175s0f0@0

# ./pktgen_sample05_flow_per_thread.sh -i enp175s0f0 \
-s 64 -d 10.0.1.1 -m 7c:fe:90:57:ab:c0 -n 0
```

and

```
./pktgen_sample05_flow_per_thread.sh -i enp175s0f0 \
-s 64 -d 10.0.1.1 -m 7c:fe:90:57:ab:c0 -n 0 -t 32
```


# kernel's pktgen (outcome)



## DPDK's pktgen (config)

```
enable 0 range
range 0 dst ip 10.0.1.2 10.0.1.2 10.0.1.254 0.0.0.1
range 0 src ip 10.0.1.3 10.0.1.3 10.0.1.254 0.0.0.1
range 0 proto udp
range 0 dst port 1 1 65534 1
range 0 src port 1 1 65534 1
range 0 dst mac 7c:fe:90:57:ab:c8 7c:fe:90:57:ab:c8
                    7c:fe:90:57:ab:c8 00:00:00:00:00:00
```

# DPDK's pktgen (outcome)



# How does the receiver feel?

```

1 [ 0.0%] 9 [|||||] 100.0% 17 [ 0.0%] 25 [|||||] 100.0%
2 [ 0.0%] 10 [|||||] 100.0% 18 [ 1.3%] 26 [|||||] 100.0%
3 [ 0.0%] 11 [|||||] 100.0% 19 [ 0.0%] 27 [|||||] 100.0%
4 [ 0.0%] 12 [|||||] 100.0% 20 [ 0.0%] 28 [|||||] 100.0%
5 [ 0.0%] 13 [|||||] 100.0% 21 [ 0.0%] 29 [|||||] 100.0%
6 [ 0.7%] 14 [|||||] 100.0% 22 [ 0.0%] 30 [|||||] 100.0%
7 [ 0.0%] 15 [|||||] 100.0% 23 [ 0.0%] 31 [|||||] 100.0%
8 [ 0.0%] 16 [|||||] 100.0% 24 [ 0.0%] 32 [|||||] 100.0%
Mem[|||||] 3.23G/93.1G Tasks: 25, 187 thr; 17 running
Swp[ ] 0K/1.50G Load average: 15.88 10.01 4.40
Uptime: 8 days, 08:52:52

```

# With iptables

```
# iptables -A INPUT -p udp -m udp -j DROP
```



# With iptables

```
# iptables -A INPUT -p udp -m udp -j DROP
```

```

1 [ 0.0%] 9 [|||||] 100.0% 17 [ 0.0%] 25 [|||||] 100.0%
2 [ 0.0%] 10 [|||||] 100.0% 18 [ 1.3%] 26 [|||||] 100.0%
3 [ 0.0%] 11 [|||||] 100.0% 19 [ 0.0%] 27 [|||||] 100.0%
4 [ 0.0%] 12 [|||||] 100.0% 20 [ 0.0%] 28 [|||||] 100.0%
5 [ 0.0%] 13 [|||||] 100.0% 21 [ 0.0%] 29 [|||||] 100.0%
6 [ 0.7%] 14 [|||||] 100.0% 22 [ 0.0%] 30 [|||||] 100.0%
7 [ 0.0%] 15 [|||||] 100.0% 23 [ 0.0%] 31 [|||||] 100.0%
8 [ 0.0%] 16 [|||||] 100.0% 24 [ 0.0%] 32 [|||||] 100.0%
Mem[|||||] 3.23G/93.1G Tasks: 25, 187 thr; 17 running
Swp[ ] 0K/1.50G Load average: 15.88 10.01 4.40
Uptime: 8 days, 08:52:52

```

# Can we do better?

# Can we do better?

```
# iptables -t raw -A PREROUTING -p udp -m udp -j DROP
```

# Can we do better?

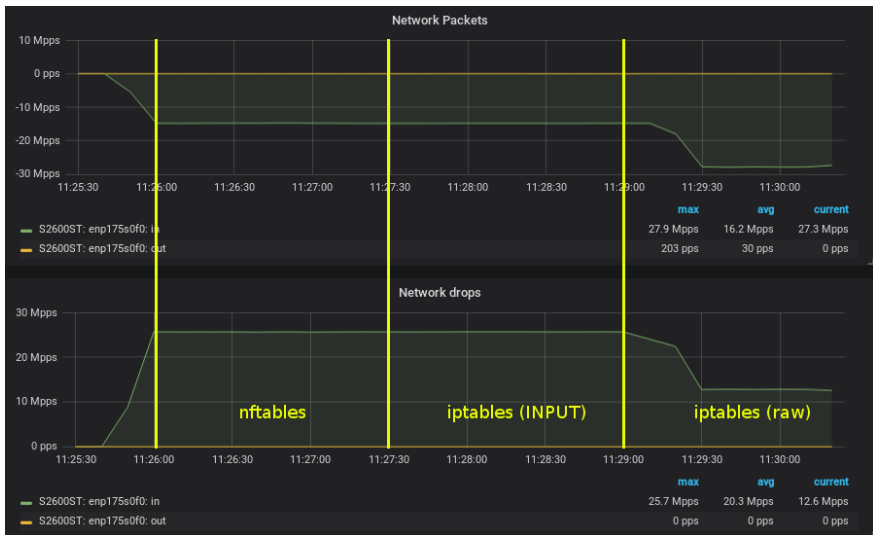
```
# iptables -t raw -A PREROUTING -p udp -m udp -j DROP
```

```

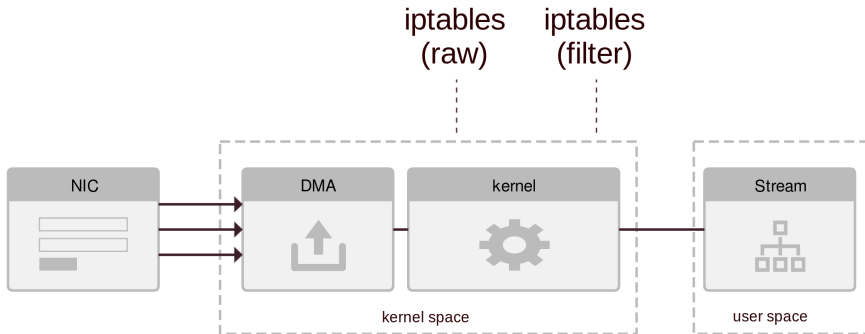
1 [ | 1.3%] 9 [ | 100.0%] 17 [ | 0.0%] 25 [ | 20.7%]
2 [ | 0.0%] 10 [ | 16.1%] 18 [ | 0.0%] 26 [ | 100.0%]
3 [ | 0.0%] 11 [ | 26.2%] 19 [ | 0.0%] 27 [ | 100.0%]
4 [ | 0.0%] 12 [ | 36.4%] 20 [ | 0.7%] 28 [ | 100.0%]
5 [ | 0.0%] 13 [ | 100.0%] 21 [ | 0.0%] 29 [ | 100.0%]
6 [ | 0.0%] 14 [ | 100.0%] 22 [ | 0.0%] 30 [ | 100.0%]
7 [ | 0.0%] 15 [ | 100.0%] 23 [ | 0.0%] 31 [ | 100.0%]
8 [ | 0.0%] 16 [ | 100.0%] 24 [ | 0.0%] 32 [ | 100.0%]
Mem [ | 1.84G/93.1G] Tasks: 30, 185 thr; 13 running
Swp [ | 0K/1.50G] Load average: 13.05 13.40 13.02
Uptime: 1 day, 19:56:50

```

# nftables and iptables



# synthesis



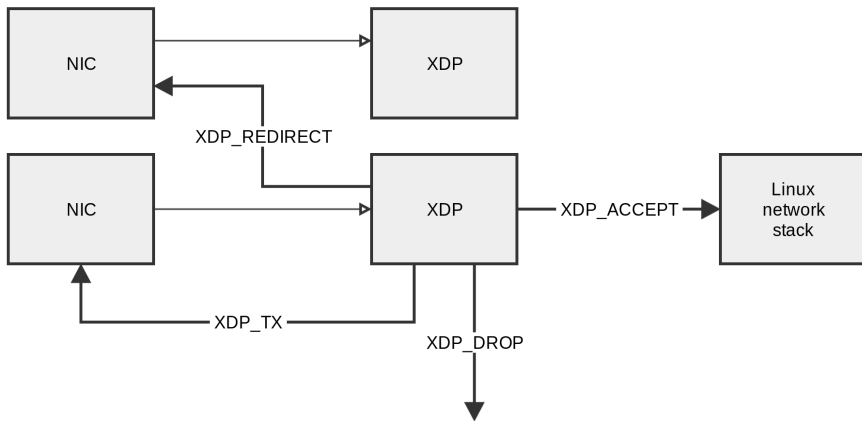
## Not the expected result

«Iptables is not slow. It's just executed too late in the stack.»

– (r) Gilberto Bertin

# Introduce XDP : What is XDP?

- ▶ XDP stands for eXpress Data Path.
- ▶ Programmable, High-performances, specialized application, packet processor in the linux networking stack.







# XDP : eXpress Data Path

- ▶ XDP is *not*:
  - ▶ a replacement for TCP/IP stack
  - ▶ kernel bypass
- ▶ Runs eBPF program on hooks:
  - ▶ In the kernel (TC/xdp-generic)
  - ▶ In driver (xdp or xdpoffload) => before **skb** allocation
- ▶ 3 outcomes:
  - ▶ Accept the packet: XDP\_PASS
  - ▶ Drop the packet: XDP\_DROP
  - ▶ Redirect the packet: XDP\_TX or XDP\_REDIRECT

# XDP



# XDP



# Minimal example

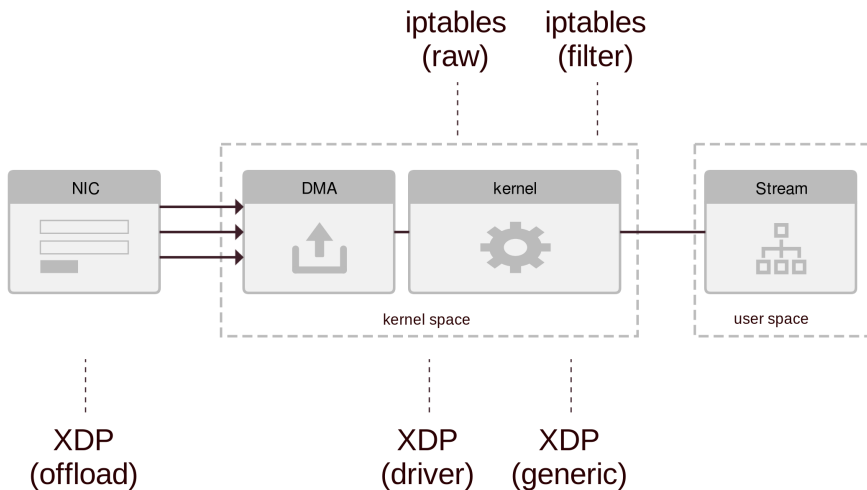
```
#include <linux/bpf.h>

#ifdef __section
# define __section(NAME) \
    __attribute__((section(NAME), used))
#endif

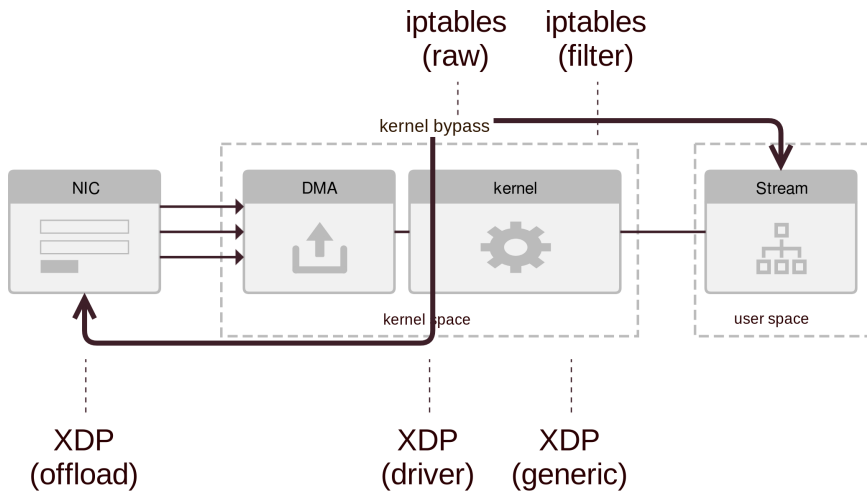
__section("prog")
int xdp_drop(struct xdp_md *ctx)
{
    return XDP_DROP;
}

char __license[] __section("license") = "GPL";
```

# Synthesis



# XDP alternatives: kernel bypass



# Kernel bypass

- ▶ PF\_RING
- ▶ NetMap
- ▶ DPDK
- ▶ ...
- ▶ Pros:
  - ▶ Fast!
- ▶ Cons:
  - ▶ Require driver support
  - ▶ Handle the whole stack “by hand”
  - ▶ NIC may be dedicated (not visible from the Linux).

## Summary and conclusion



# What we have seen

- ▶ Scaling traffic is not trivial;
- ▶ Filters need to be applied as early as possible;
- ▶ XDP is a standard (as in mainline integrated) way;
- ▶ But alternatives exist.

# Issues with XDP

- ▶ Require “recent” software stack
  - ▶ kernel
  - ▶ iproute
  - ▶ toolchain (LLVM for instance)
- ▶ Complex
  - ▶ Basically have to know C
- ▶ Increasing number of tools
  - ▶ bpfILTER
  - ▶ bcc
  - ▶ P4

# Try it yourself

Fork me on github : <https://github.com/fser/pts-2018>

# References

- ▶ <https://jvns.ca/blog/2017/04/07/xdp-bpf-tutorial/>
- ▶ <https://qmonnet.github.io/whirl-offload/2016/09/01/dive-into-bpf/>
- ▶ <https://cilium.readthedocs.io/en/latest/bpf/>
- ▶ <https://www.iovisor.org/technology/xdp>
- ▶ <http://prototype-kernel.readthedocs.io/en/latest/bpf/index.html>
- ▶ man pages:
  - ▶ tc-bpf (8)
  - ▶ man bpf (2)
- ▶ Documentation/networking/filter.txt
- ▶ Several netdev-conference's slides.

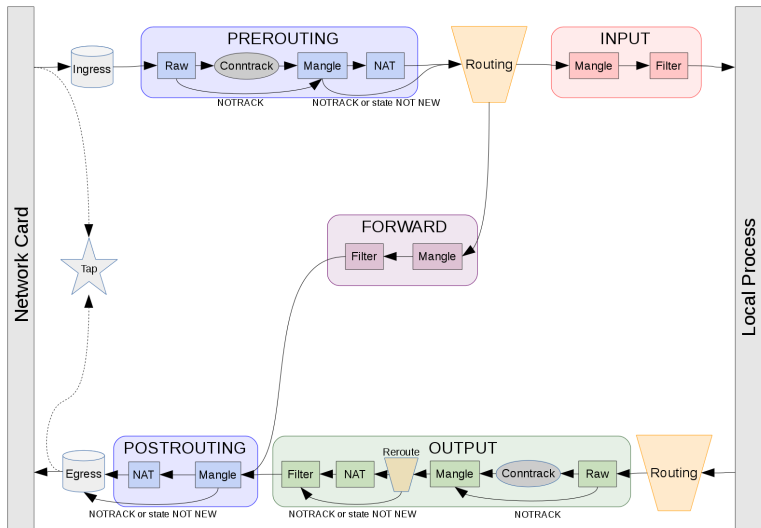
# Questions

## Backup slides

# Loading an XDP program

```
# ip link set dev DEVICE xdp \  
    obj OBJECT_FILE.o [ sec SECTION_NAME ]  
  
# tc qdisc add dev DEVICE clsact  
# tc filter add dev DEVICE ingress bpf da obj OBJECT_FILE.o
```

# Iptables overview





# Flood memcached commands

```
#!/usr/bin/env python
```

```
import sys, socket, random
```

```
UDP_IP, UDP_PORT = "127.0.0.1", 11211
```

```
MESSAGE = "\x00\x00\x00\x00\x00\x01\x00\x00{}\r\n"
```

```
cmds = '''add append cas decr delete flush_all  
get gets incr prepend replace stats'''.split()
```

```
sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
while True:
```

```
    cmd = random.choice(cmds)
```

```
    sock.sendto(MESSAGE.format(cmd), (UDP_IP, UDP_PORT))
```

# XDP parsing - bcc

```
#!/usr/bin/env python
```

```
from bcc import BPF
```

```
...
```

```
b = BPF(src_file="xdp_memcached.c", cflags=["-w",  
      "-DRETURNCODE=%s" % ret, "-DCTXTYPE=%s" % ctxtype])
```

```
b.attach_xdp(device, fn, flags)
```

```
dropcnt = b.get_table("dropcnt")
```

# Licenses

Memcached traffic viewer: Apache License, Version 2.0

XDP UDP drop: GPL v2

Scripts & ansible: WTFPL

Slides

